



# Kubernetes and Container Management

## **Container Camp 2015**

Tim Hockin <[thockin@google.com](mailto:thockin@google.com)>  
Senior Staff Software Engineer  
[@thockin](#)



Google Cloud Platform





Google has been developing  
and using containers to  
manage our applications for  
**over 10 years.**



# Everything at Google runs in containers:

- Gmail, Web Search, Maps, ...
- MapReduce, batch, ...
- GFS, Colossus, ...
- Even **GCE itself**: VMs run in containers



Shipping Containers At Clyde, by Steve Gibson

# Everything at Google runs in containers:

- Gmail, Web Search, Maps, ...
- MapReduce, batch, ...
- GFS, Colossus, ...
- Even **GCE itself**: VMs run in containers

## We launch over **2 billion** containers **per week**.



Shipping Containers At Clyde, by Steve Gibson



# But it's so different!

- Deployment
- Management, monitoring
- Isolation (very complicated!)
- Updates
- Discovery
- Scaling, replication, sets

A **fundamentally different** way of managing **applications** requires different tooling and abstractions



# Enter Kubernetes

Greek for *“Helmsman”*; also the root of the word *“Governor”* and *“cybernetic”*

- Container orchestrator
- Runs and manages containers
- Supports multiple cloud and bare-metal environments
- Inspired and informed by Google’s experiences and internal systems
- **100% Open source**, written in **Go**

Manage applications, not machines





# Pets vs. Cattle



# A 50000 foot view




# A 50000 foot view

users

master





nodes



# A 50000 foot view




Run X  
Replicas = 2  
Memory = 4Gi  
CPU = 2.5






# A 50000 foot view




# A 50000 foot view







# A 50000 foot view




# A 50000 foot view







# A 50000 foot view



# A 50000 foot view




# A 50000 foot view





# Container clusters: A story in two parts

# All you really care about



# Container clusters: A story in two parts

## 1. Setting up a cluster

- Choose a cloud: GCE, AWS, Azure, Rackspace, **on-premises**, ...
- Choose a node OS: CoreOS, Atomic, RHEL, Debian, CentOS, Ubuntu, ...
- Provision machines: Boot VMs, install and run kube components, ...
- Configure networking: IP ranges for Pods, Services, SDN, ...
- Start cluster services: DNS, logging, monitoring, ...
- Manage nodes: kernel upgrades, OS updates, hardware failures...

Not the easy or fun part, but unavoidable

This is where things like **Google Container Engine (GKE)** really help



# Container clusters: A story in two parts

## 2. Using a cluster

- Run Pods & Containers
- Replication controllers
- Services
- Volumes
- Secrets

This is the fun part!


A distinct set of problems from cluster setup and management

Don't make developers deal with cluster administration!

Accelerate development by focusing on the applications, not the cluster

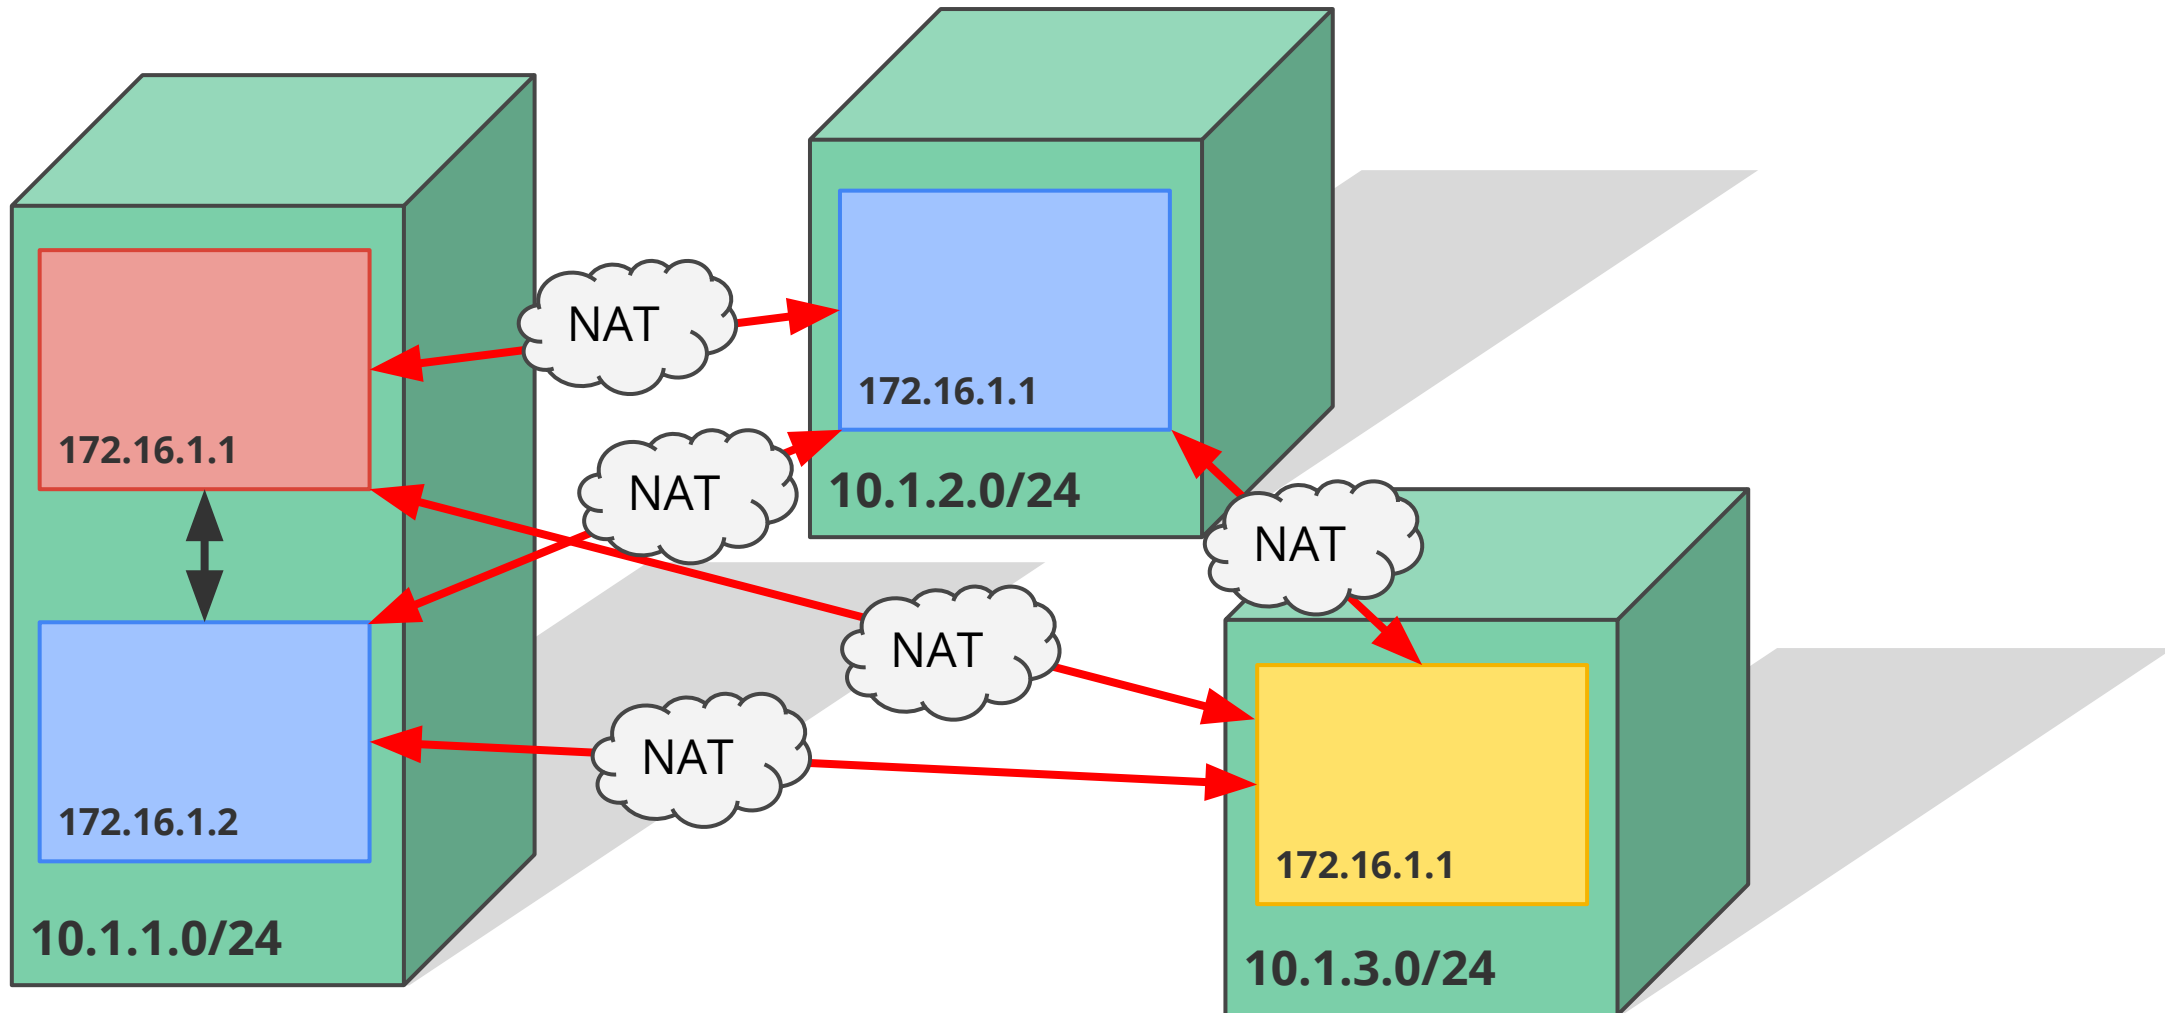
# Networking

# Docker networking





# Docker networking



# Kubernetes networking

## Pod IPs are **routable**

- docker default is private IP

## Pods can reach each other without NAT


- even across nodes

## No **brokering** of port numbers


- too complex, why bother?

## This is a **fundamental requirement**

- several SDN solutions exist



# Kubernetes networking





# Concept: Pods

**Small group** of containers & volumes

**Tightly** coupled

The atom of scheduling & placement in Kubernetes


Shared namespace

- share IP address & localhost
- share IPC

Mortal

- can die, cannot be reborn

**Example: data puller & web server**



# Concept: Services

A group of pods that **work together**

- grouped by a selector

Defines access policy

- “load balanced” or “headless” for now


Gets a stable **virtual IP** and port

- called the service *portal*
- also a DNS name

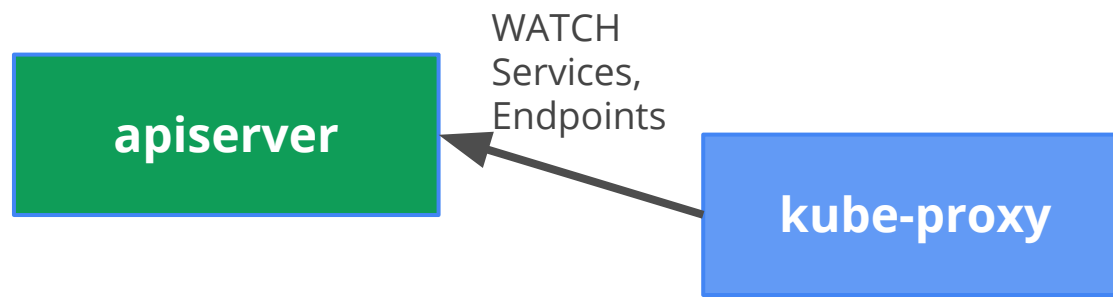
VIP is captured by *kube-proxy*

- watches the service constituency
- updates when backends change

Hides complexity - ideal for non-native apps



# Services



# Services

## Pod

- Name = "pod1"
- Labels = {"App": "Nifty"}
- Port = 9376

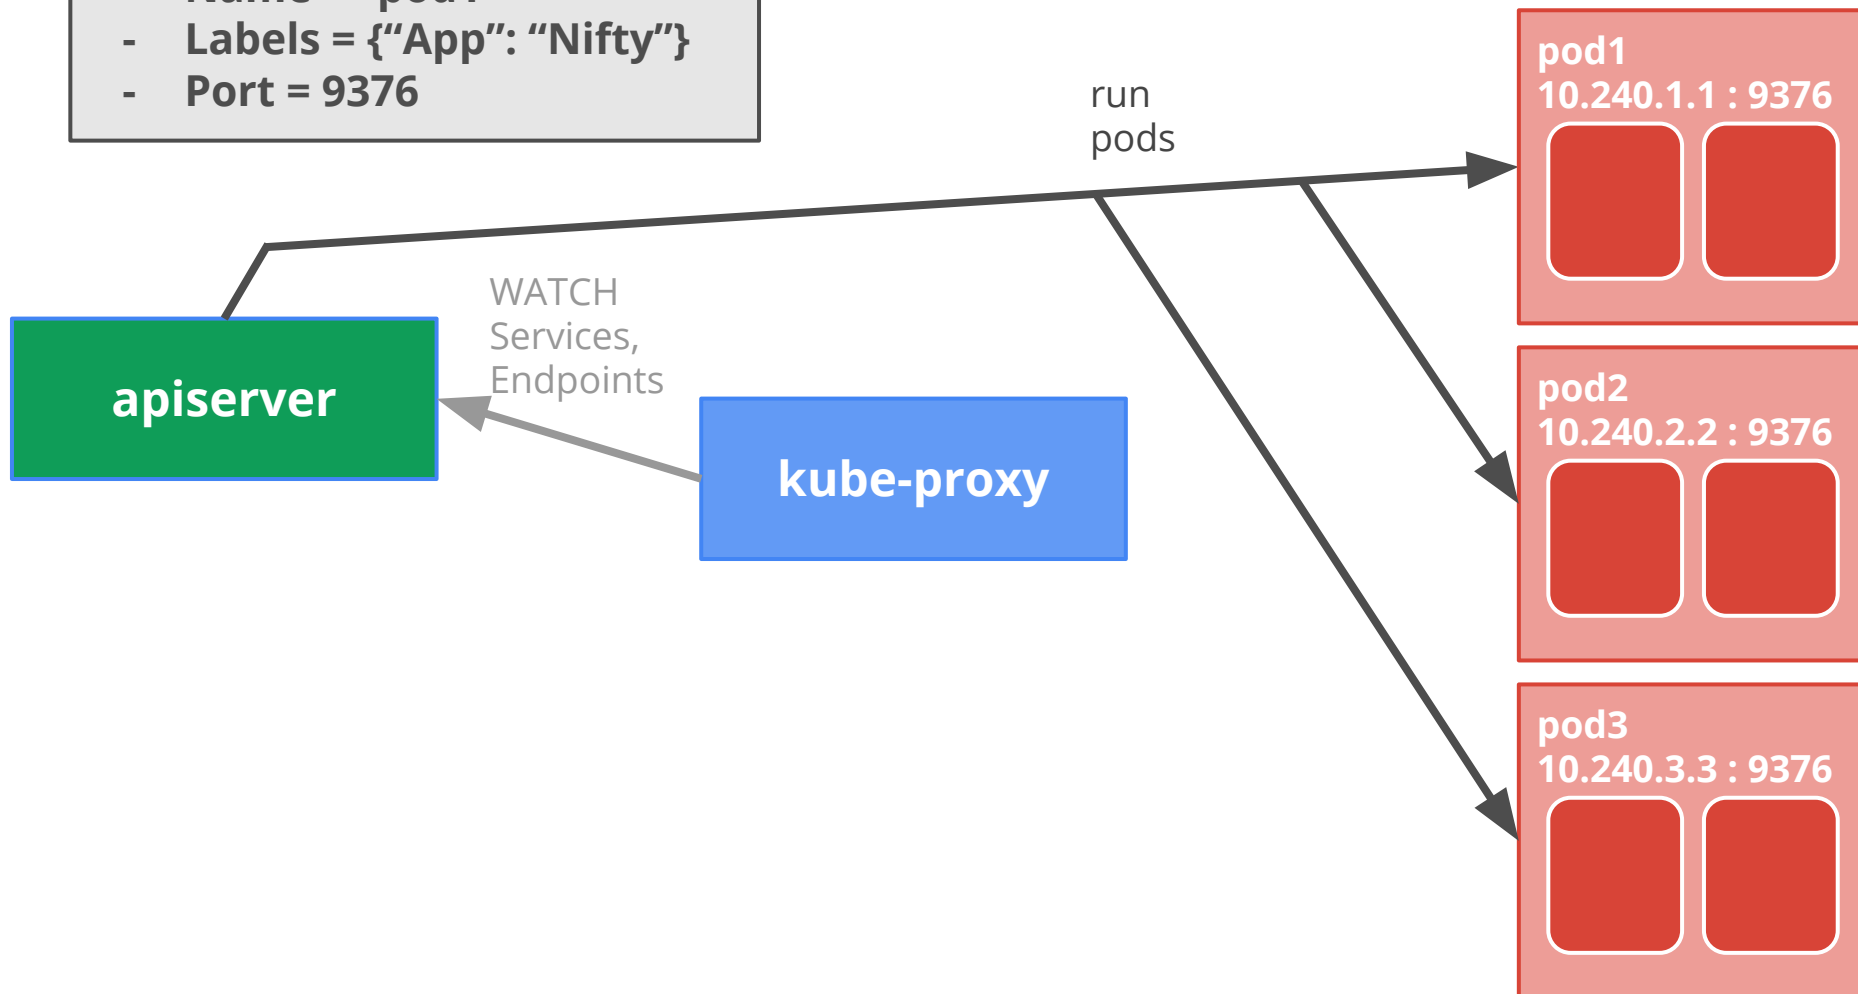




# Services

## Pod

- Name = "pod1"
- Labels = {"App": "Nifty"}
- Port = 9376



# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



pod1  
10.240.1.1 : 9376



pod2  
10.240.2.2 : 9376




pod3  
10.240.3.3 : 9376



# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



pod1  
10.240.1.1 : 9376



pod2  
10.240.2.2 : 9376




pod3  
10.240.3.3 : 9376



# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



pod1  
10.240.1.1 : 9376



pod2  
10.240.2.2 : 9376



pod3  
10.240.3.3 : 9376

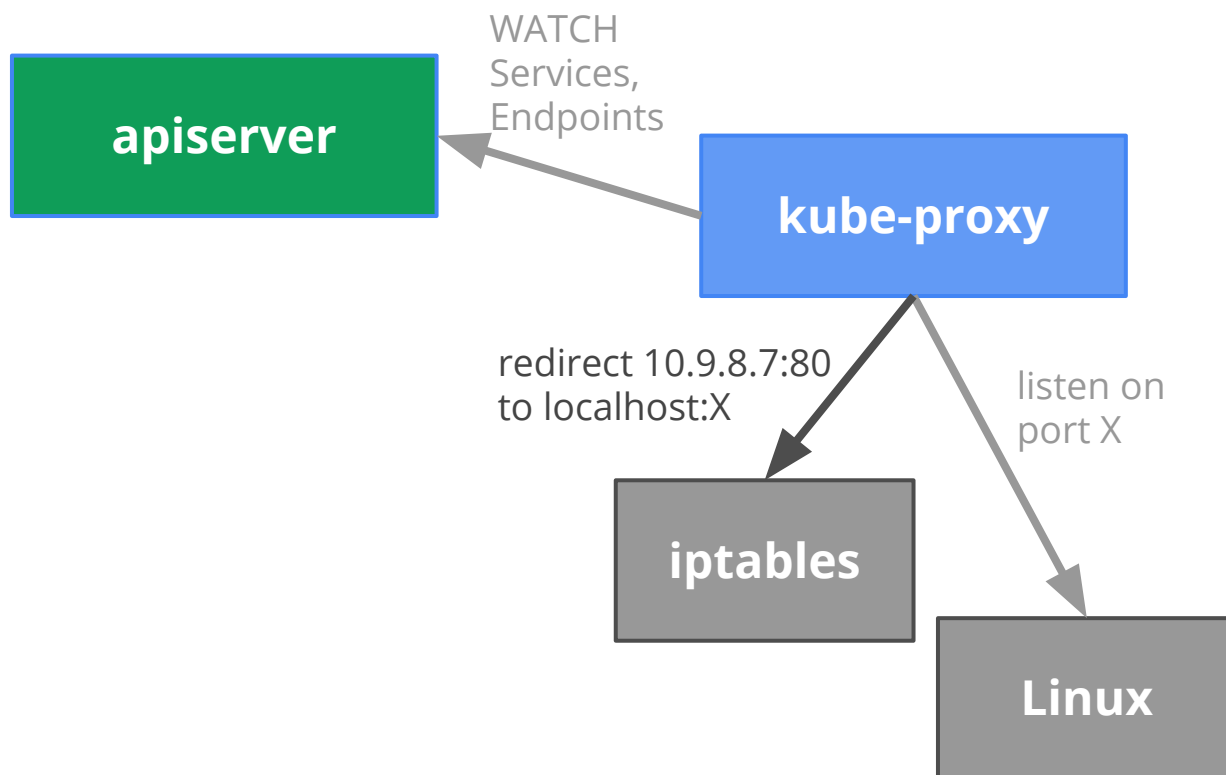




# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



pod1  
10.240.1.1 : 9376



pod2  
10.240.2.2 : 9376



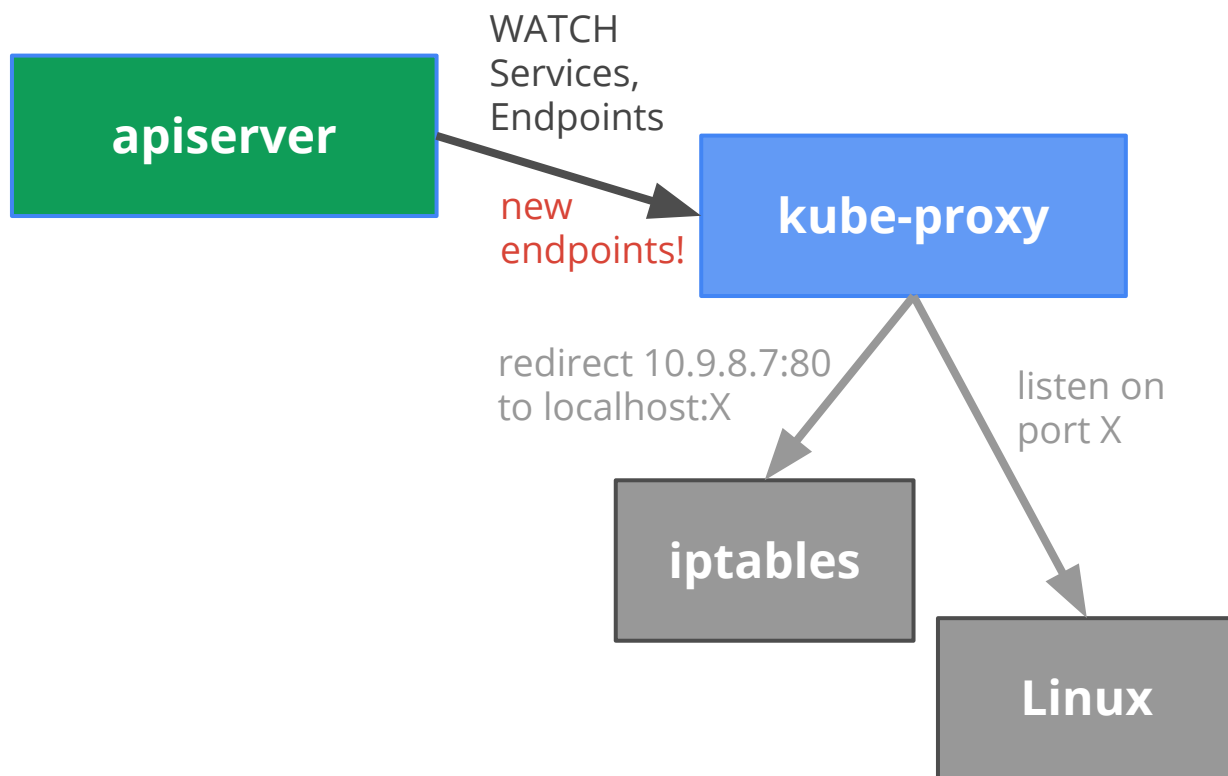
pod3  
10.240.3.3 : 9376



# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



pod1  
10.240.1.1 : 9376



pod2  
10.240.2.2 : 9376



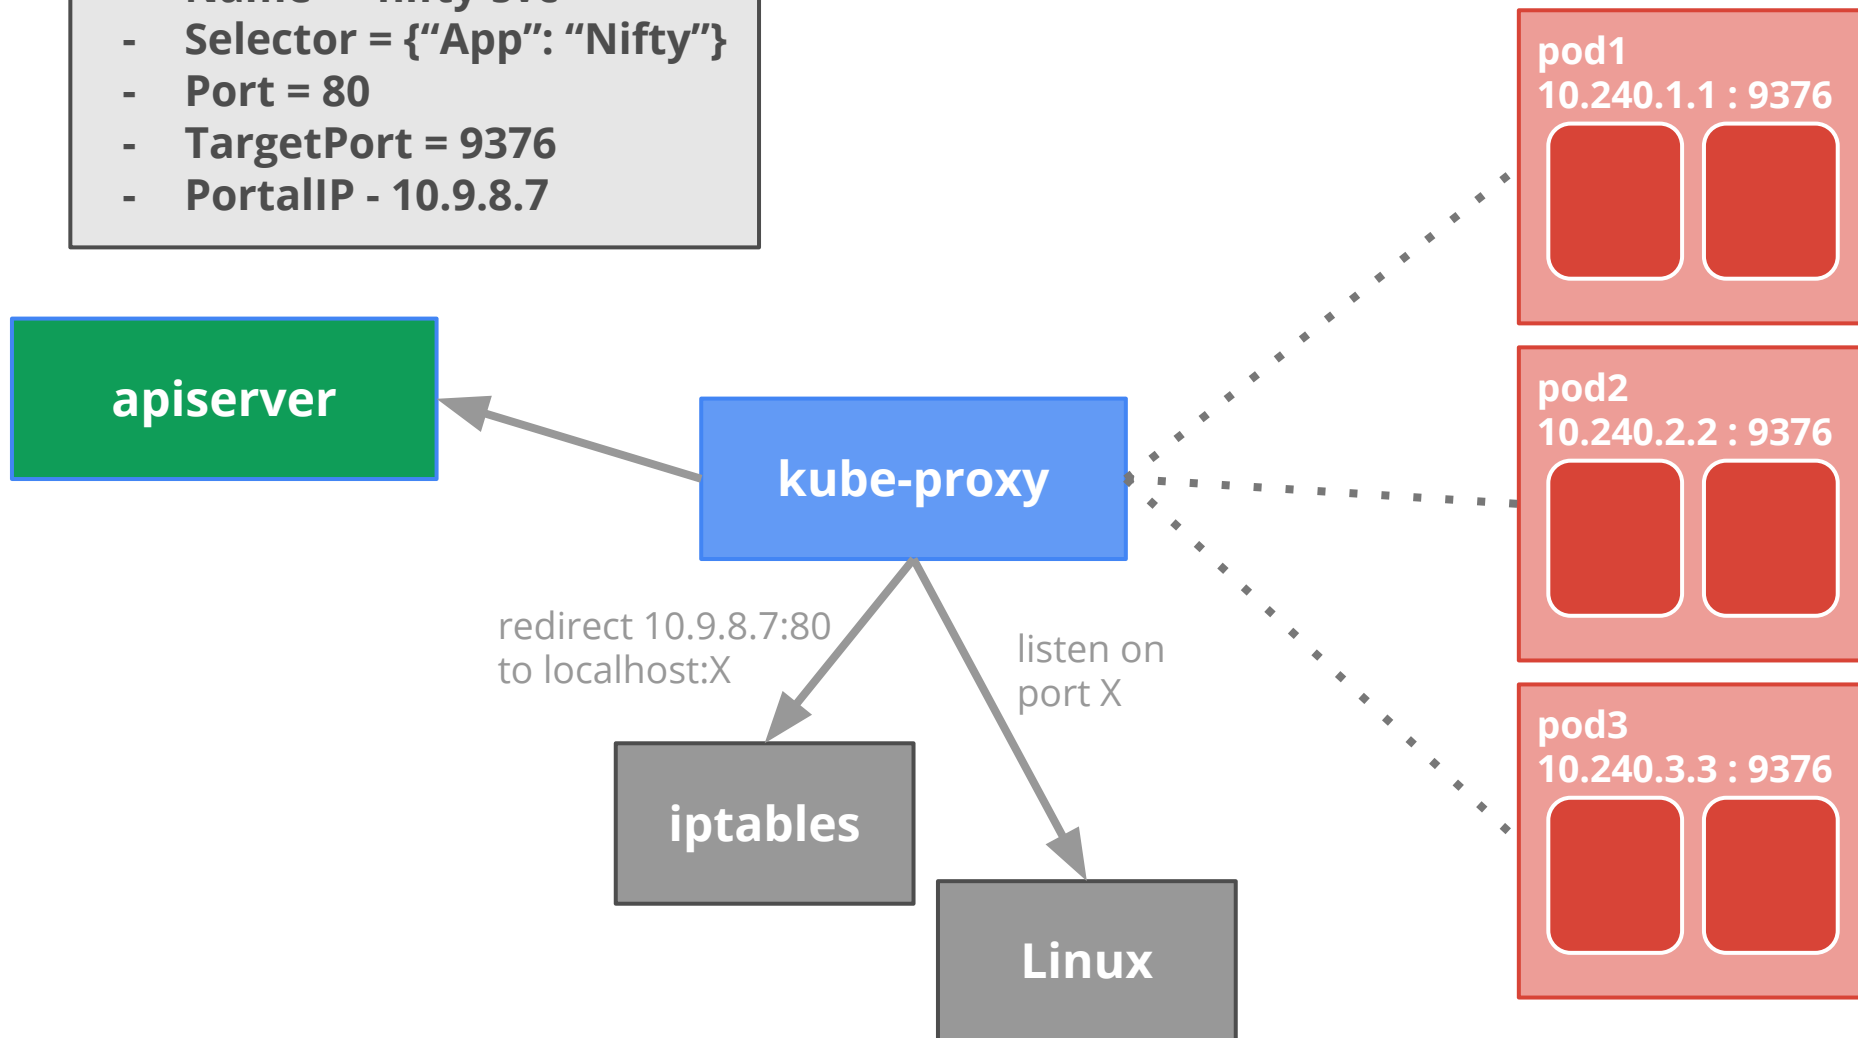
pod3  
10.240.3.3 : 9376



# Services

## Service

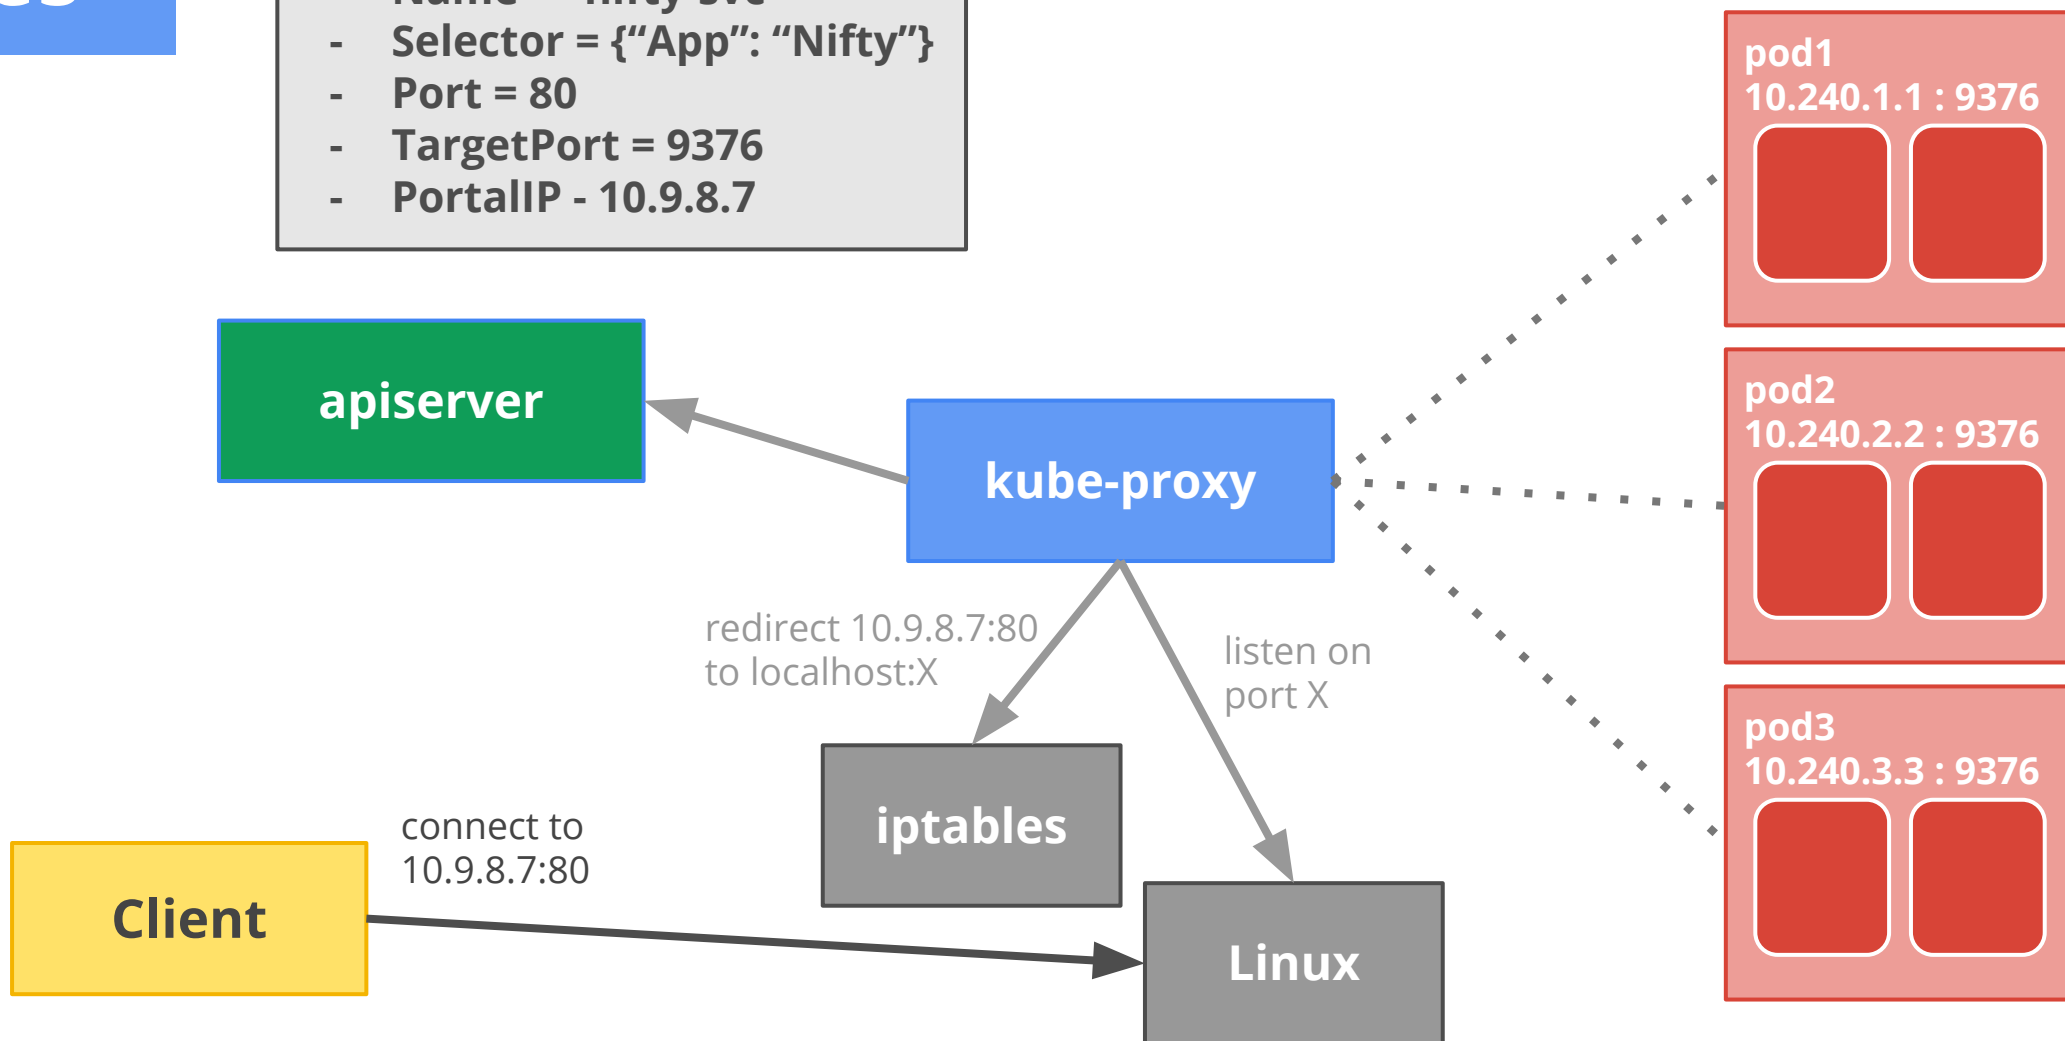
- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



# Services

## Service

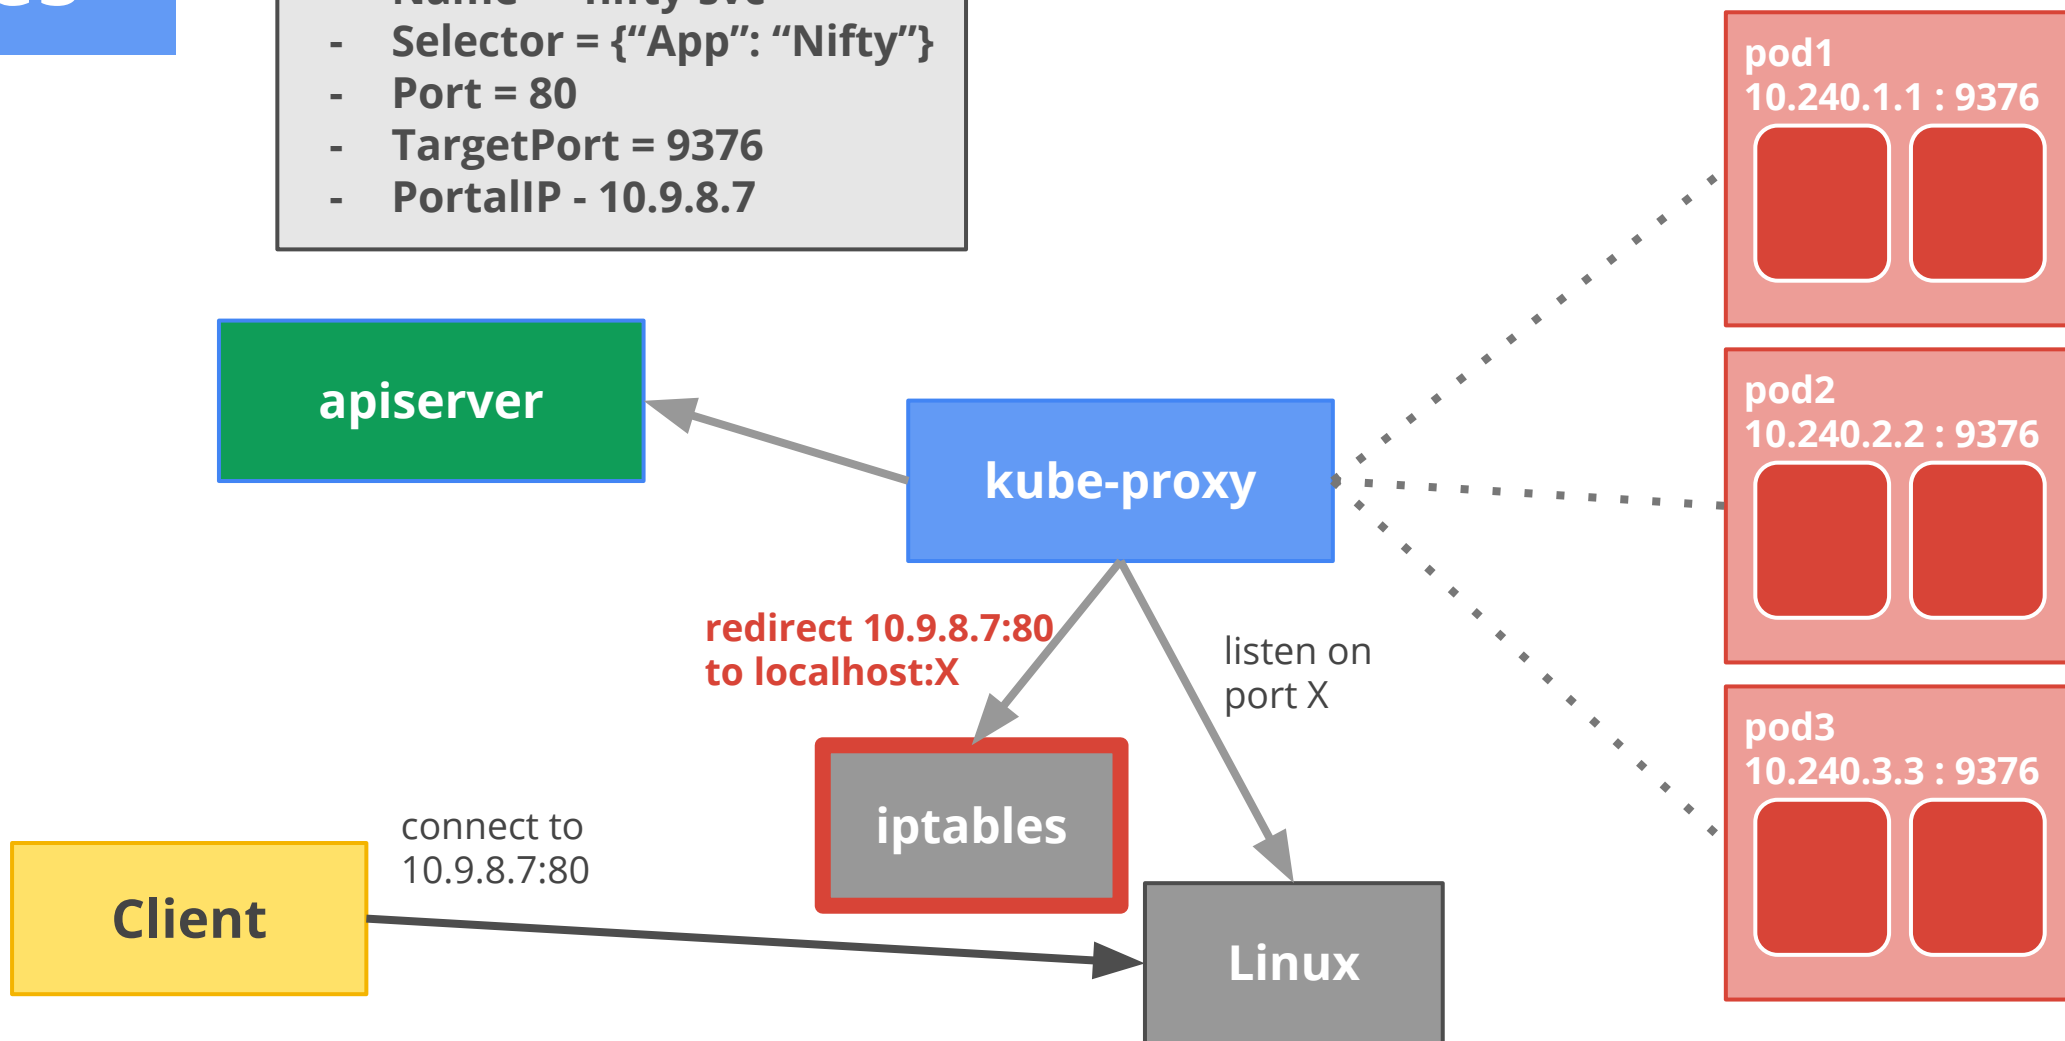
- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



# Services

## Service

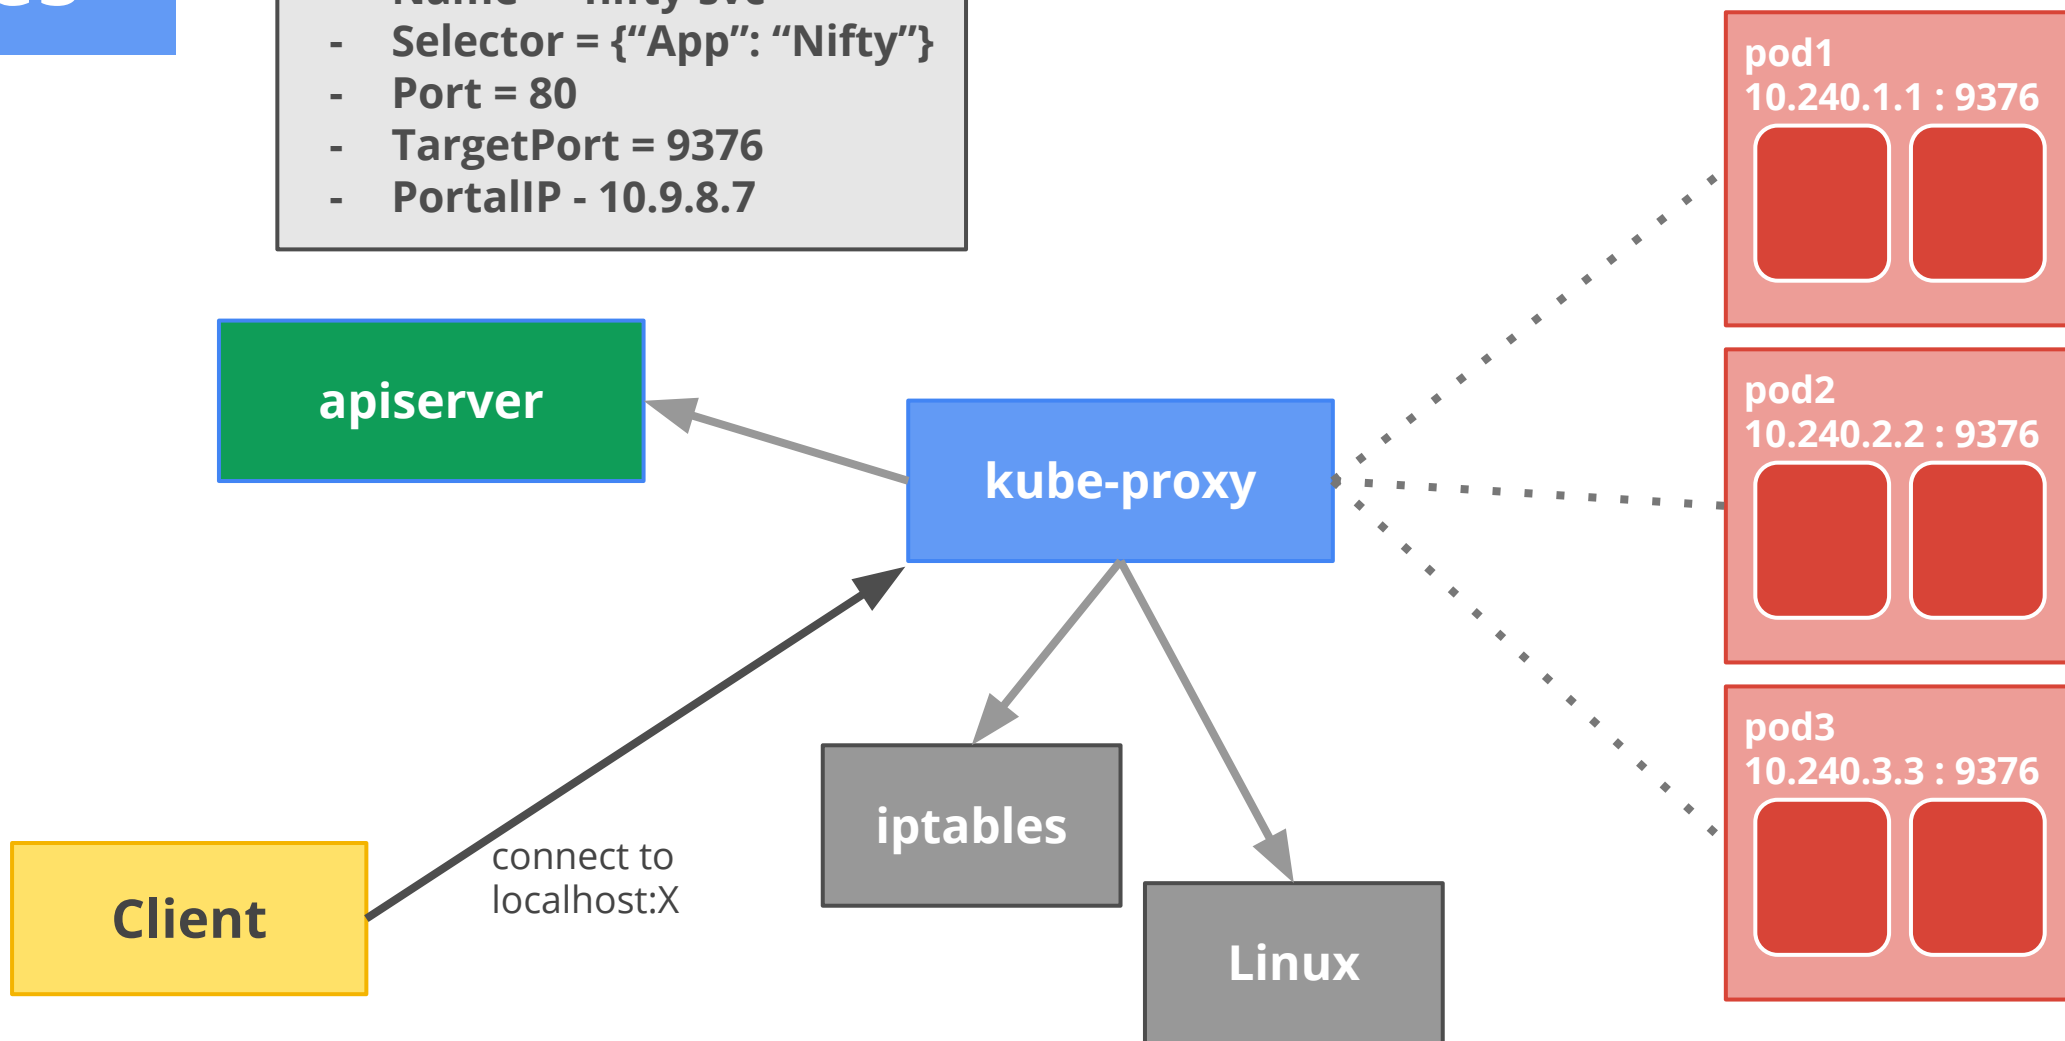
- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7

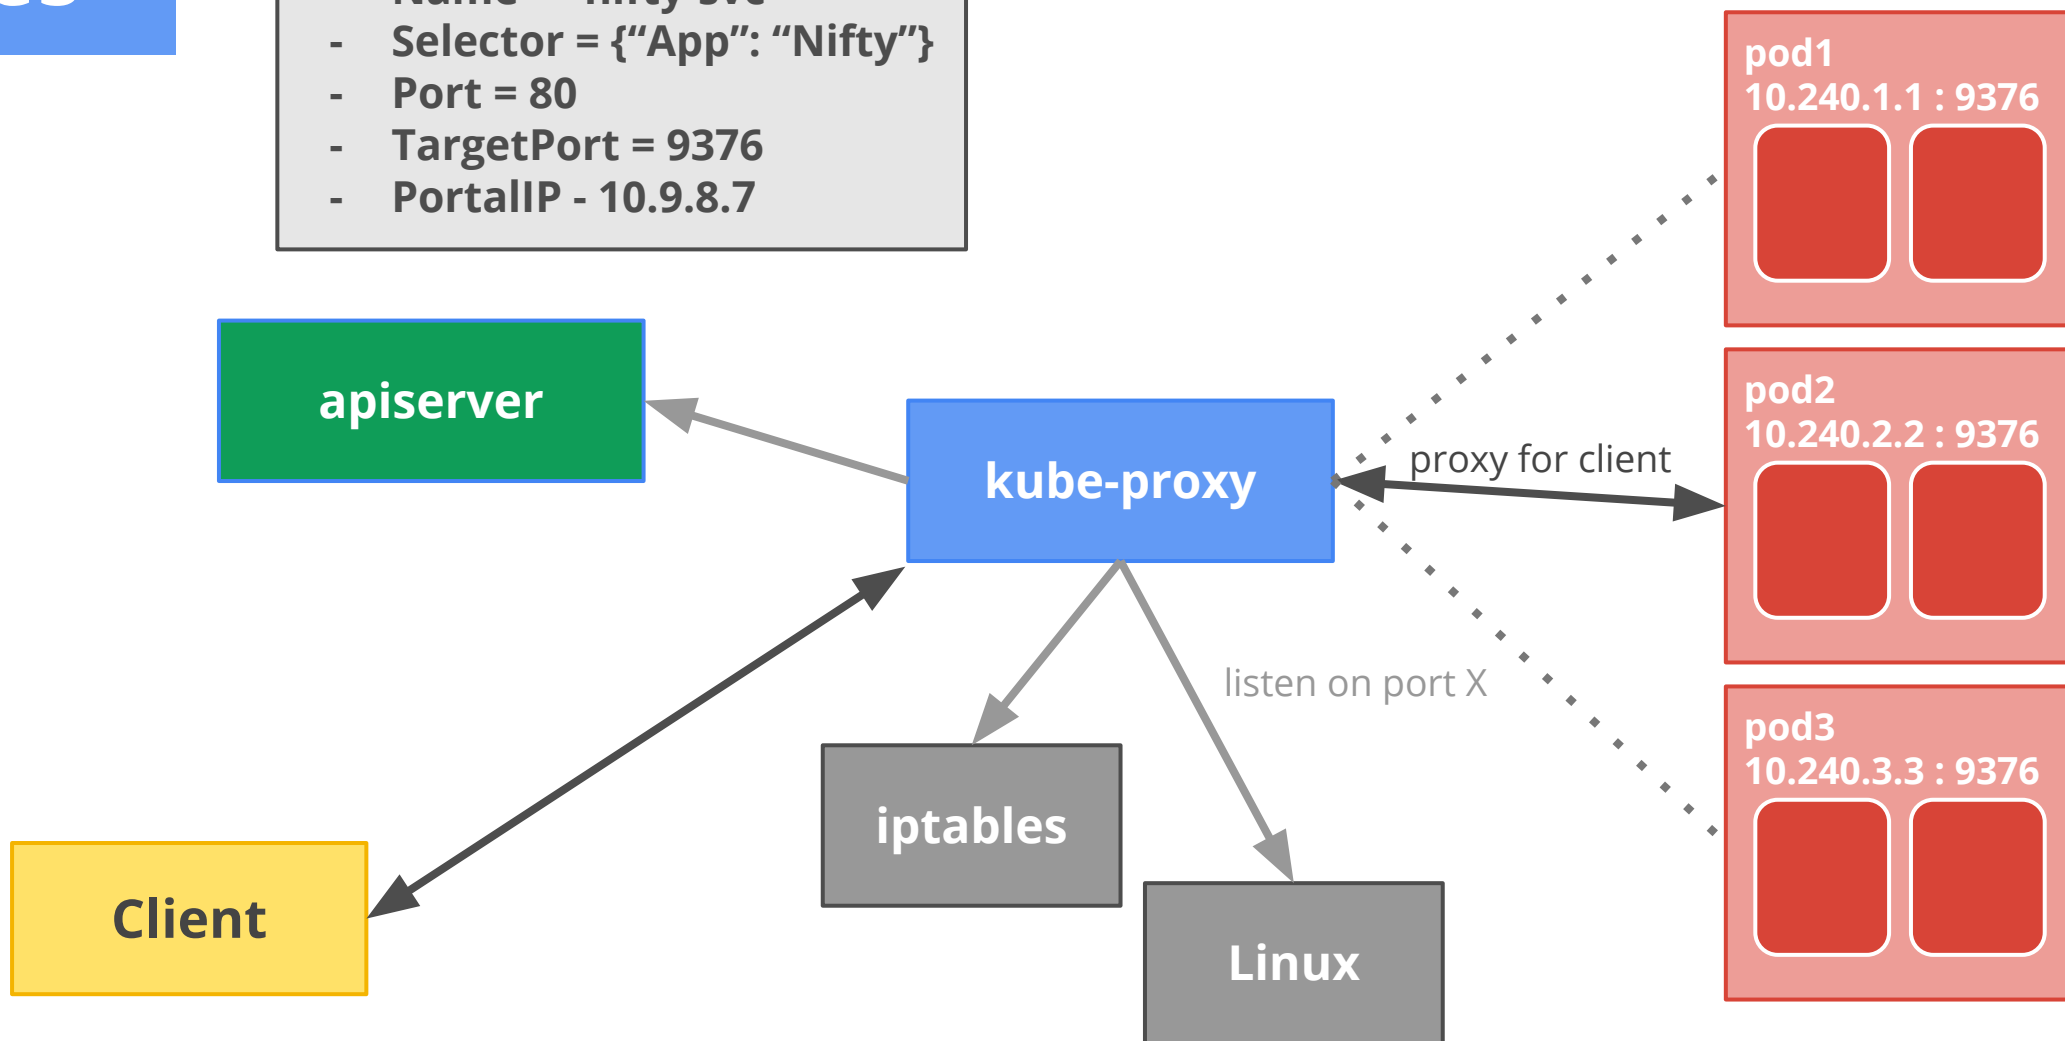




# Services

## Service

- Name = "nifty-svc"
- Selector = {"App": "Nifty"}
- Port = 80
- TargetPort = 9376
- PortalIP - 10.9.8.7



# Concept: Volumes


Very similar to Docker's concept

Pod scoped storage

Share the pod's lifetime & fate

Support many types of volume plugins

- Empty directory
- Host path
- Git repository
- GCE Persistent Disk
- AWS Elastic Block Store
- iSCSI
- NFS
- GlusterFS
- Ceph File and RBD
- Cinder
- ...



# New: Persistent Volumes

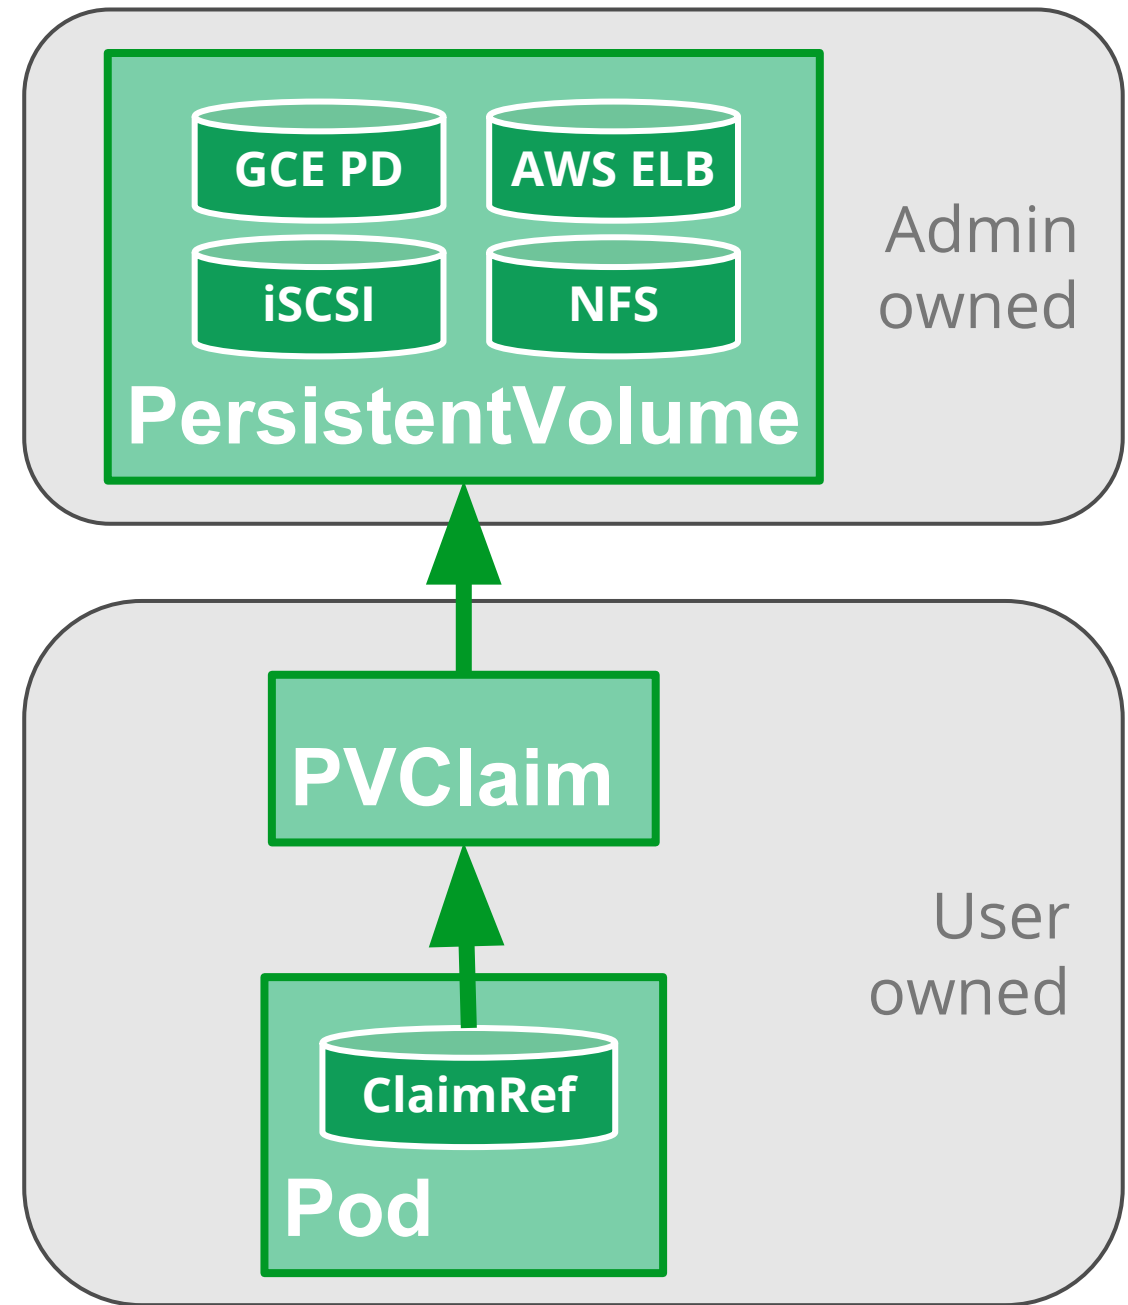
A higher-level abstraction - insulation from any one cloud environment

Admin provisions them, users claim them

Independent lifetime and fate

Can be handed-off between pods and lives until user is done with it

Dynamically “scheduled” and managed, like nodes and pods



# Docker, Rocket, LXC, Oh my!


Currently built on Docker

Work is in progress to abstract that (a bit) into a Runtime abstraction

Interest in Rocket and LXC support

Rocket support is in flight (we like plugins)

Dynamically “scheduled” and managed, like nodes and pods



# What else is in new?

- Network plugins
- Secrets
- Graceful termination
- Quota
- More volumes
- Downward API
- More platforms
- Performance
- Scalability
- High availability masters
- Scheduling
- Cluster federation
- Multi-cloud
- Easier setup

# Kubernetes status & plans

Open sourced in June, 2014

- won the 2014 BlackDuck “rookie of the year” award

Google Container Engine (GKE)

- hosted Kubernetes - don't think about cluster setup

Red Hat: OpenShift 3

- open PaaS on Kubernetes

CoreOS: Tectonic

- ready-to-run Kubernetes - don't think about cluster setup

Mirantis: Murano

- Kubernetes and OpenStack

Driving towards a 1.0 release in O(months)

Roadmap:

- <https://github.com/GoogleCloudPlatform/kubernetes/blob/master/docs/roadmap.md>





# The Goal: Shake things up

Containers are a **new way of working**

Requires new concepts and new tools

Google has a **lot** of experience...

...but we are **listening to the users**

**Workload portability is important!**



# Kubernetes is Open

- open community
- open design
- open to ideas
- open source

<http://kubernetes.io>

<https://github.com/GoogleCloudPlatform/kubernetes>

irc.freenode.net *#google-containers*  
@kubernetesio